

Advanced Programming Techniques Sheet 1 — Tensor Class

This sheet is not mandatory. However, it teaches essential concepts about C++ programming and also guides you how the core data structure (i.e. the tensor class) of the project can be implemented. Thus, we strongly recommend you work on all these exercises.

Part 1 — Basic C++ Quizzes

Complete the following Studon tests under "Exercise Material -> StudOn Tests":

a) Function resolution - Template function specialization Demonstrate your even more advanced knowledge in the resolution of functions by taking a quick StudOn Quiz ("Function resolution - templates").

b) Function resolution - Mixed concepts Show off your extreme knowledge in the resolution of functions by taking a quick StudOn Quiz ("Function resolution - challenge").

Part 2 — Multidimensional arrays / Tensors

In this exercise, you will write a simple multidimensional array data structure. Its *shape* is a tuple of positive integers that defines how many elements are stored in each direction. The resulting implementation shall resemble a simplified version of [numpy's ndarray](#).

Multidimensional arrays are crucial for computations in numerical linear algebra. In particular they are the workhorse data structure in software that is used to approximate partial differential equations and their solutions as well as in machine learning applications. The latter usually uses the term *tensor* to refer to the same data structure¹. Clearly, every engineer or researcher should be able to write their tools from scratch before using them in production. So we will do that now.

Before working on this exercise, you should download and extract the corresponding files from StudOn.

Compile the exercise with CMake and Make. We ship the exercises with a main function and tests. Eventually you can run the tests to verify your implementation. We are using **Wall**, **Wextra**, **Werror** and **pedantic** flags. These flags should enforce you to write clean code.

Thorough out the exercise, there will be links to other resources and to the [C++ Core Guidelines](#). Please follow the guidelines, as they immensely help with writing solid code.

a) Tensor class-template Implement your very own Tensor class-template, including the following features:

- arbitrary shape (with an arbitrary number of dimensions),
- arbitrary component type (to instantiate integer tensors, floating-point tensors, etc.),
- copy- and move-constructors and -assignment operators,
- element accessors,
- (de-)serialization (reading and writing tensors to files).

¹Have a look at the tensor classes in [TensorFlow](#) and [PyTorch](#).

A skeleton is given in `tensor.hpp`. Regrettably, the exercise might be ahead of the lecture on some topics. However, we think that teaching you modern C++ with the correct habits is more important than convenience. (We ignore guideline [C.100](#) for this educational task.)

Subtasks

- *Structure.* Think about how you want to implement the `Tensor` class-template, especially what member variables your class should have. Add these member variables in `tensor.hpp`. You can add (private) helper methods if you want to. You are only allowed to change `tensor.hpp`. The changes must work with the other provided files.
 - [C.9](#)
 - [C.12](#)
- *Concepts.* The class template is associated with a constraint using the concepts feature introduced in C++20 (see [cppreference](#)). Tensors shall only be allowed to hold arithmetic types (e.g. `int`, `float`).
 - [T.concepts](#)
- *Constructors.* Implement a primary constructor that receives the shape of the tensor. The shape defines the number of elements in each dimension. An empty shape vector results in a scalar, a shape of size 1 is equivalent to a vector with `shape[0]` elements. A shape of size 2 results in a tensor that resembles a matrix. The number of dimensions (or size of the shape-vector) is also referred to as *rank*. The default constructor shall create a tensor with rank 0, containing a single, zero-initialized entry. Use `std::fill` to fill up the memory with the `fillValue` if specified.
 - [C.40](#)
 - [C.41](#)
- *Copy-constructor.* Implement the copy-constructor, that takes a reference to an existing `Tensor`, and returns a `Tensor` with the same shape and contents. Use `std::copy` for copying the contents.
 - [C.21](#)
 - [C.61](#)
 - [C.101](#)
- *Move-constructor.* Implement the move-constructor, too. The move-constructor behaves similarly to the copy-constructor, except that it may reuse the storage of the `Tensor` whose rvalue reference we receive.

`std::move` is C++'s answer on having very stable resource management while retaining high performance. It is crucial that you understand the difference between `lvalues` and `rvalues`. [Move Semantics](#) is a tremendous introduction to this topic.

 - [C.64](#)
 - [C.66](#)
 - [C.102](#)

`std::exchange` makes a highly readable code.

After moving, the moved rvalue shall be equal to a default-constructed tensor.
- *Destructor.* Implement the destructor. Make sure to free any resources that you allocated in the constructor.
 - [C.30](#)

- C.31
- C.32
- C.33
- C.36
- *Assignment.* Implement the assignment operator and a move assignment operator.
 - C.62
 - C.65
- *Accessors.* Implement `operator()` for read and write access to the tensor elements. The element index is passed as a vector of size `rank`. Indexing starts at 0.
- *Member functions.* Implement the `rank`, `shape`, and `numElements` member functions.
 - C.4
- *File IO.* Implement the functions `writeTensorToFile` and `readTensorFromFile`. The file format is defined as follows:

```

<rank>
<shape[0]>
...
<shape[rank-1]>
<t(0, 0, ..., 0, 0)>
<t(0, 0, ..., 0, 1)>
...
<t(0, 0, ..., 0, shape[rank-1] - 1)>
<t(0, 0, ..., 1, 0)>
<t(0, 0, ..., 1, 1)>
...
<t(shape[0] - 1, shape[1] - 1, ..., shape[rank-1] - 1)>

```

First the `rank` (i.e. size of the shape array) is specified. Next all `rank-1` entries of the shape array. Eventually all elements of the tensor follow. All entries are separated by spaces.

- *Testing.* Make sure that all the tests in `test_tensor.cpp` work with your implementation.
- *Memory leaks.* Use `valgrind`² to check for memory leaks.
- *Readability.* Check your code. Can you enhance the readability?

b) Matrix-vector multiplication Many applications frequently use matrices and vectors. So it would be convenient to specialize our tensor class-template for that purpose. Use the tensor class-template to write your own matrix and vector class-templates. The skeletons are supplied in `matvec.hpp`. This should be quite straightforward as many features can be reused from the tensor implementation. Then implement matrix-vector multiplication and make sure that all tests in `test_matvec.cpp` pass.

²`valgrind ./test_tensor -leak-check=full`